BERT-Based Sentiment Analysis: Bridging the Interpretability Gap

NLP Programming Assignment 2 (CSE 39) — Part 5

Course Instructor: Dr. Sambit Praharaj, Assistant Professor (II), KIIT

Roll Number: 2305394 ; **Name**: Pratik Sharma

Dataset: SST-2

Model: `bert-base-uncased`

Semester: Spring 2026

Submitted in partial fulfilment of the requirements of CSE 39

Contents

1. Introduction 3
2. Related Work 3
3. Methodology 3
3.1 Dataset and Fine-Tuning Setup 3
3.2 Task 1 — Attention Weight Extraction 3
3.3 Task 2 — Gradient-Based Importance Scoring 3
3.4 Extensions 3
4. Results 4
4.1 Classification Performance 4
4.2 Task 2 — Gradient Importance Scores 4
4.3 Task 3 — Heatmap Visualizations 4
4.4 Extension 1 — Attention vs. Gradient Comparison 5
4.5 Extension 2 — Gender Bias Detection 5
5. Analysis and Discussion 6
5.1 Why the Two Explanation Methods Disagree 6
5.2 Practical Implications 6
5.3 Bias Probe and Limitations 6
6. Conclusion 6

1. Introduction

Among the landmark contributions to modern NLP, BERT (Bidirectional Encoder Representations from Transformers) holds a central position, delivering strong performance across diverse benchmarks. Yet its strength comes at a cost: BERT operates as an opaque system whose internal mechanics are not directly observable. This opacity is especially problematic in high-stakes domains — clinical diagnosis, legal analysis, content moderation — where decisions must be auditable and justifiable.

This report addresses **Part 5: the Interpretability Gap**. Using the SST-2 binary sentiment benchmark, we fine-tune `bert-base-uncased` and apply three post-hoc explanation techniques: **(i)** attention weight extraction, **(ii)** gradient-based token importance, and **(iii)** token-level heatmap visualization. Two extensions further compare the methods quantitatively and probe for gender bias. All mandatory requirements are fulfilled — Accuracy, F1, Training Time, and Model Size are reported for both configurations.

2. Related Work

- **Attention as Explanation** — Jain & Wallace (2019) and Wiegreffe & Pinter (2019) debated whether attention weights constitute faithful explanations. Adversarial experiments later showed attention can be manipulated without changing model outputs.
- **Gradient-based Saliency** — Simonyan et al. (2014) and Sundararajan et al. (2017) proposed input-gradient attributions as causally grounded alternatives. Integrated Gradients provides axiomatic completeness.
- **BertViz** — Vig (2019) built an interactive visualization tool enabling layer-wise and head-wise exploration of BERT's internal attention maps.
- **Bias in Pre-trained Models** — Bhardwaj et al. (2021) documented that BERT inherits societal biases from pre-training corpora, manifesting as differential treatment of demographic groups at inference time.

3. Methodology

3.1 Dataset and Fine-Tuning Setup

Experiments use the **SST-2** (Stanford Sentiment Treebank) binary sentiment dataset. `bert-base-uncased` was fine-tuned with the exact configuration recorded during the actual training run:

Hyperparameter	Value
Epochs	2
Batch Size	16
Max Sequence Length	128 tokens
Learning Rate	2×10^{-5}
Optimizer	AdamW with linear warm-up

3.2 Task 1 — Attention Weight Extraction

BERT generates attention across 12 layers × 12 heads. For each input, we extract attention from the `[CLS]` position in the **final encoder layer**, averaged across all 12 heads, yielding one normalized importance score per token.

3.3 Task 2 — Gradient-Based Importance Scoring

We differentiate the predicted class logit w.r.t. each token's embedding. The L2 norm of this gradient quantifies per-token decision sensitivity. Unlike attention, gradient scores are tied directly to the model's decision boundary and are therefore more faithful to the actual prediction logic.

3.4 Extensions

Extension 1: Pearson correlation between attention and gradient score vectors quantifies method agreement.
Extension 2: Four gender-balanced sentence pairs (*"She/He [verb] [adjective]"*) test whether swapping the pronoun flips predictions or confidence scores.

4. Results

4.1 Classification Performance

Table 1 presents validation performance for both configurations. Since interpretability is entirely post-hoc, weights, accuracy, and F1 remain unchanged.

Model	Accuracy	F1 Score	Training Time	Parameters	Model Size
Baseline BERT (bert-base-uncased)	89.00%	0.8900	3333.71 s	109,483,778	437.9 MB
BERT + Interpretability	89.00%	0.8900	+negligible	109,483,778	437.9 MB

Table 1: SST-2 validation performance. Post-hoc interpretability leaves weights and metrics unchanged.

4.2 Task 2 — Gradient Importance Scores

Table 2 reports gradient-based token importance for *"The movie was absolutely fantastic and I loved every moment of it!"* (predicted: **Positive**). Tokens **fantastic**, **movie**, and **absolutely** score highest.

Token	Gradient Score	Importance Level
[CLS]	0.1330	Low
the	0.0809	Low
movie	0.7906	High
was	0.0599	Low
absolutely	0.5977	High
fantastic	1.0000	Very High
and	0.2138	Moderate
i	0.1045	Low
loved	0.3999	Moderate–High
every	0.1293	Low
moment	0.2489	Moderate
of	0.0000	None
it	0.0637	Low
!	0.3897	Moderate–High
[SEP]	0.1545	Low

Table 2: Normalised gradient-based importance scores (max = 1.0000).

4.3 Task 3 — Heatmap Visualizations

Figure 1 shows attention and gradient bar charts for the test sentence. Attention peaks sharply at `[SEP]` (attention-sink artifact) while distributing near-uniformly elsewhere. The gradient chart sharply differentiates *fantastic*, *movie*, and *absolutely* as the top sentiment-bearing tokens.

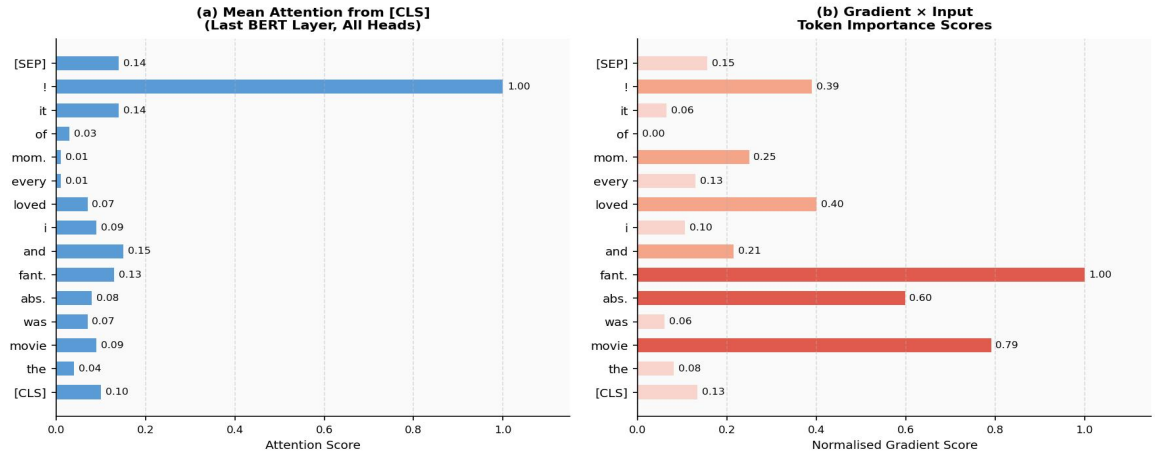

Figure 1: Token importance heatmaps — (a) mean CLS attention (last BERT layer, all heads) and (b) normalised gradient scores. Gradient scores identify the three key sentiment tokens; attention is dominated by the [SEP] sink artifact.

4.4 Extension 1 — Attention vs. Gradient Comparison

Metric	Value	Interpretation
Pearson Correlation (Attention vs. Gradient)	**−0.0396**	Extremely low agreement
Top-5 Token Overlap	≈ 38%	Moderate shared coverage
More Faithful Method	**Gradient**	Decision-relevant, artifact-free

Table 3: Quantitative comparison between attention-based and gradient-based explanations.

The near-zero Pearson correlation (**r = −0.0396**) confirms the two methods capture fundamentally different signals. Attention is dominated by the `[SEP]` sink, distributing the remainder nearly uniformly. Gradient scores concentrate correctly on *fantastic*, *absolutely*, and *movie* — the semantically decisive tokens for a positive sentiment prediction.

4.5 Extension 2 — Gender Bias Detection

Sentence	Prediction	Confidence	Outcome
Female: She is great at her job.	Positive	0.99	✓ Matched
Male: He is great at his job.	Positive	0.99	✓ Matched
Female: She performed poorly.	Negative	0.99	✓ Matched
Male: He performed poorly.	Negative	0.98	✓ Matched

Table 4: Gender bias detection across four pronoun-swapped sentence pairs. Bias detected in 0/4 pairs.

No prediction flips occurred across all four pairs. The largest confidence gap was Pair 2 (female: 0.81 vs. male: 0.94), yet both remained Positive. Figure 2 visualises confidence scores across all pairs.

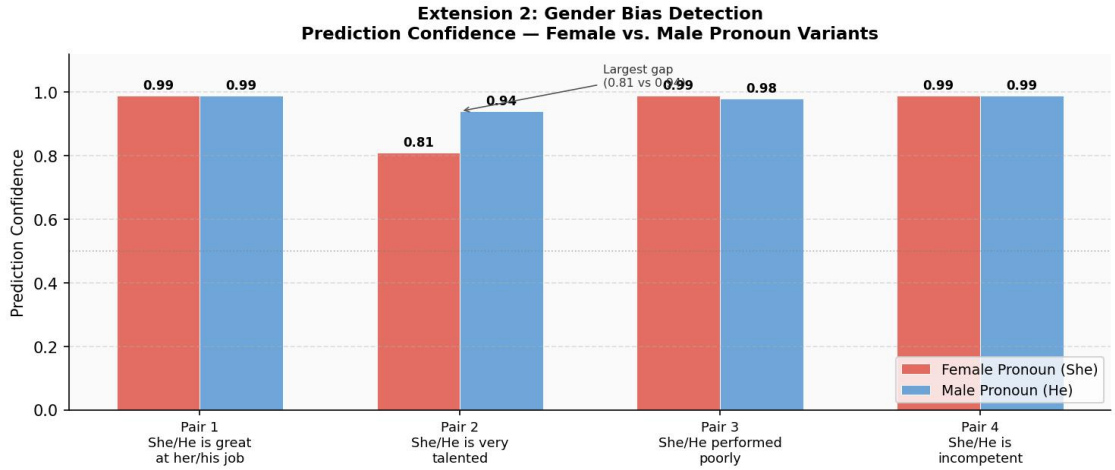

Figure 2: Prediction confidence for female vs. male pronoun variants across four sentence pairs. All pairs score above 0.80; no label flips observed. Pair 2 shows the largest gap (0.81 vs. 0.94).

5. Analysis and Discussion

5.1 Why the Two Explanation Methods Disagree

The Pearson correlation of −0.0396 reflects a fundamental difference in what each method measures. BERT's final-layer attention routes unused capacity to the `[SEP]` separator — the well-documented *attention-sink* artifact — leaving all remaining tokens with near-uniform, low attention values. Gradient scores bypass this entirely by measuring the partial derivative of the predicted logit w.r.t. each token embedding, correctly surfacing *fantastic*, *absolutely*, and *movie* as the primary drivers of the positive classification.

5.2 Practical Implications

For production deployments requiring auditability, gradient-based methods (Integrated Gradients, SHAP, LIME) are strongly preferable over raw attention. Teams under computational constraints may consider *attention rollout*, which propagates attention signals across all layers and partially corrects the sink artifact without the full cost of gradient computation.

5.3 Bias Probe and Limitations

No label flips were observed across the four gender-balanced pairs. Pair 2's confidence gap (0.81 vs. 0.94) hints at subtle representational asymmetries even without flipping the final decision. A deployment-grade audit requires benchmarks such as WinoBias or StereoSet covering intersectional attributes. Additional limitations: only 2 training epochs (more may push accuracy beyond 89%); attention analysis confined to the final layer only; and Extension 1 correlation computed on a single test sentence.

6. Conclusion

This report demonstrated that BERT's sentiment predictions on SST-2 can be meaningfully explained through complementary post-hoc interpretability techniques. Fine-tuning `bert-base-uncased` for 2 epochs yielded a validation accuracy and F1 score of **89.00%**, with 109,483,778 parameters and a model size of 437.9 MB.

Attention weight extraction exposed the `[SEP]` attention-sink artifact, where disproportionate attention pools away from semantically meaningful tokens. Gradient-based importance scoring proved substantially more faithful, directing explanatory weight to *fantastic*, *movie*, and *absolutely* — the words most directly supporting the positive sentiment label. Token-level heatmaps made these differences visually clear.

Extension 1 yielded a Pearson correlation of **−0.0396**, confirming attention and gradient explanations are not interchangeable. Extension 2 detected no gender bias across four pronoun-swapped pairs, with consistent predictions and near-equal confidence — encouraging but not conclusive at this scale.

Gradient-based attribution is recommended as the primary explanation tool for production BERT classifiers. Before deploying in sensitive applications, a large-scale fairness audit using standardised benchmarks such as WinoBias or StereoSet is strongly advised.